

Lesson 3: Model-based learning in a two-step task


Model-based planning



Model-based planning



Model-based planning






Latent learning




Edward C. Tolman


GOAL




Tolman et al., 1946 I



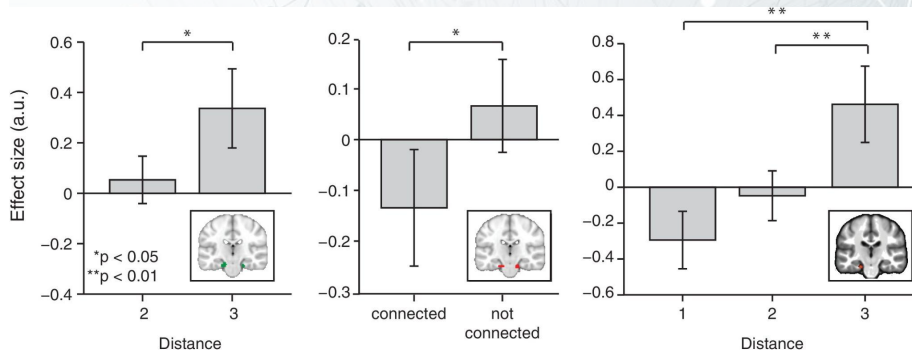

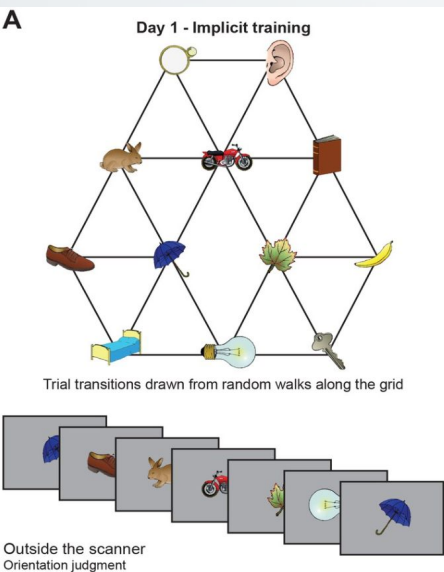
When the environment changes, model-based planning allows quick adaptation



Cognitive maps and the brain




Abstract cognitive maps





Model-based learning

(Two-step task with probabilistic transitions)

Two-step task



Two-step task



Hybrid model

Parameters

Learning rate - α

Inverse temperature - β

Eligibility factor - λ

Model based weighting - ω


Hybrid model

Model

$State = \{A, B, C\}$

$Action = \{1, 2\}$

A



B/C

Hybrid model

Model

$$State = \{A, B, C\}$$


$$Action = \{1, 2\}$$

$$Q^{MF}(s, a)$$

$$Q^{MB}(s, a)$$

A


B/C



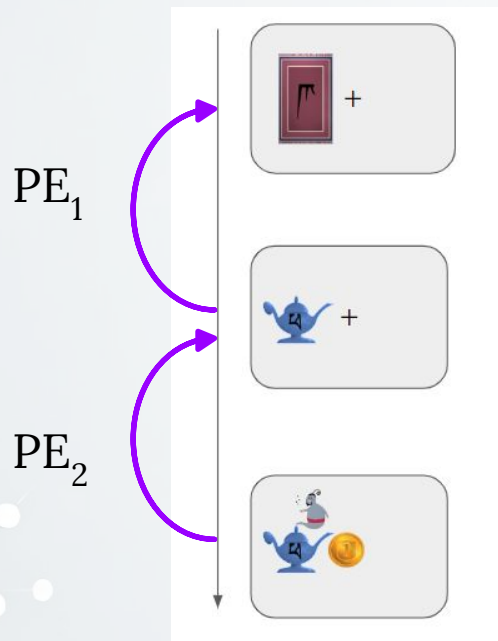
Hybrid model

A

B/C



Model-free action value updating



$$PE_2 = r_t - Q(s_2, a_1)_t$$

$$PE_1 = Q^{MF}(s_2, a_2)_t - Q^{MF}(s_1, a_1)_t$$

$$Q^{MF}(s_1, a_1)_{t+1} = Q^{MF}(s_1, a_1)_t + \alpha \cdot PE_1 + \alpha \cdot \lambda \cdot PE_2$$

$$Q^{MF}(s_2, a_2)_{t+1} = Q^{MF}(s_2, a_2)_t + \alpha \cdot PE_2$$

Hybrid model

MB action value updating

$$Q^{MB}(s_A, a_j) = P(s_B | s_A, a_j) \cdot \max_{a \in \{a_A, a_B\}} Q^{MF}(s_B, a_i) + P(s_C | s_A, a_j) \cdot \max_{a \in \{a_A, a_B\}} Q^{MF}(s_C, a_i)$$

$$P(s_B | s_A, a_1) = 0.7$$

$$P(s_B | s_A, a_2) = 0.3$$

$$P(s_C | s_A, a_1) = 0.3$$

$$P(s_C | s_A, a_2) = 0.7$$

$$Q^{MF}(s=B) = [0.5, 0.8]$$


$$Q^{MF}(s=C) = [0.2, 0.4]$$

$$Q^{MB}(s=A, a=1) = 0.7 \cdot \max[0.5, 0.8] + 0.3 \cdot \max[0.2, 0.4]$$


$$Q^{MB}(s=A, a=2) = 0.7 \cdot \max[0.2, 0.4] + 0.3 \cdot \max[0.5, 0.8]$$

A

B/C



Hybrid model



Action selection

$$Q^{\text{net}}(a) = \omega \cdot Q^{MB}(s = A, a) + (1 - \omega) \cdot Q^{MF}(s = A, a)$$

$$P(a \mid s = A) = \frac{\exp[\beta \cdot Q_{\text{net}}(a)]}{\sum_{a'} \exp[\beta \cdot Q_{\text{net}}(a')]}$$

$$P(a \mid s_2) = \frac{\exp[\beta \cdot Q^{\text{MF}}(s_2, a)]}{\sum_{a'} \exp[\beta \cdot Q^{\text{MF}}(s_2, a')]}$$

Step-by-step simulation

[google sheet](#)



Simulation code in github







Hard

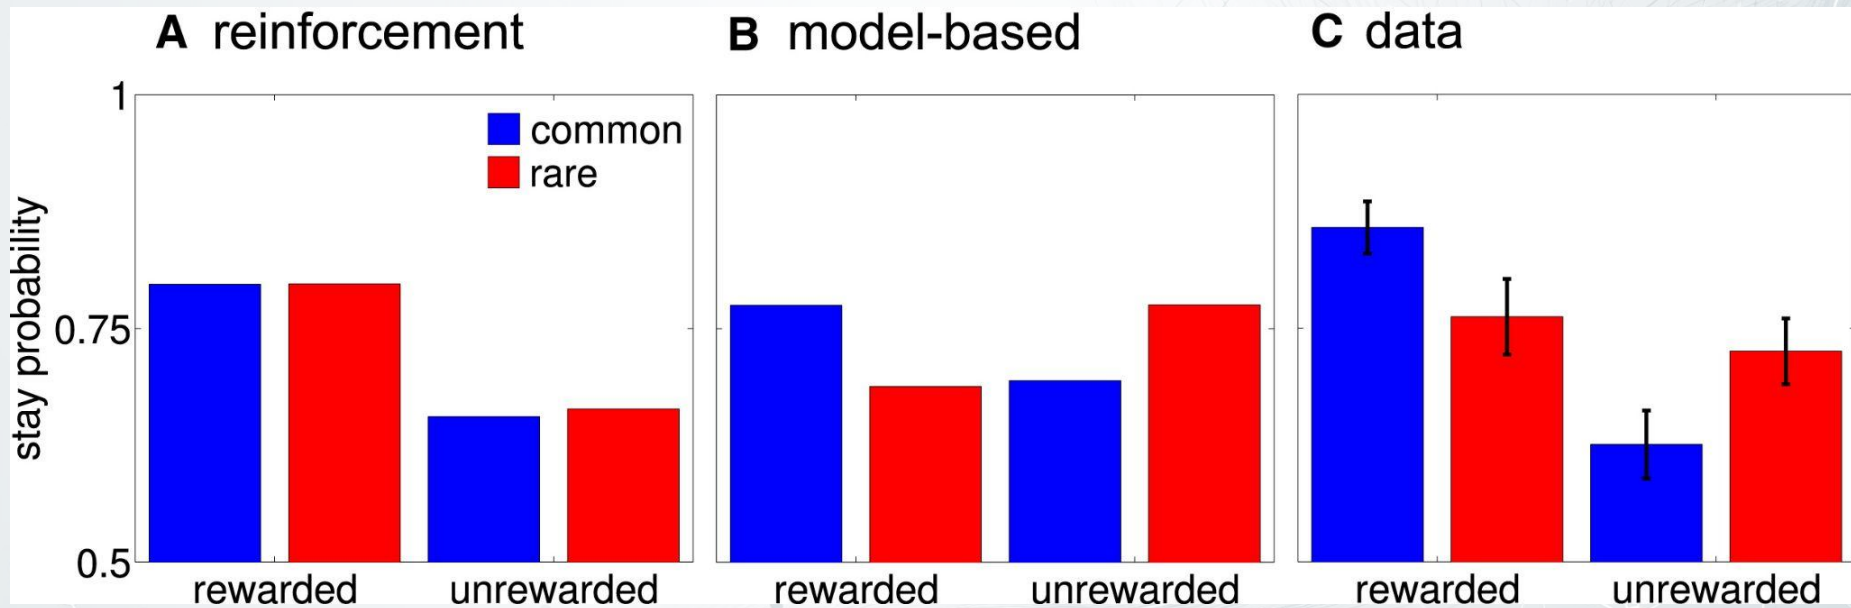
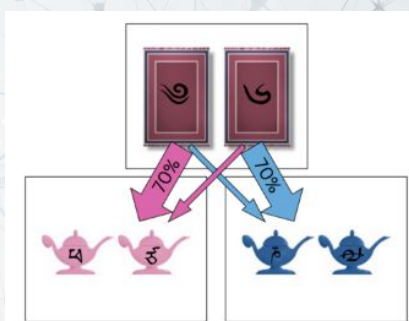


Medium

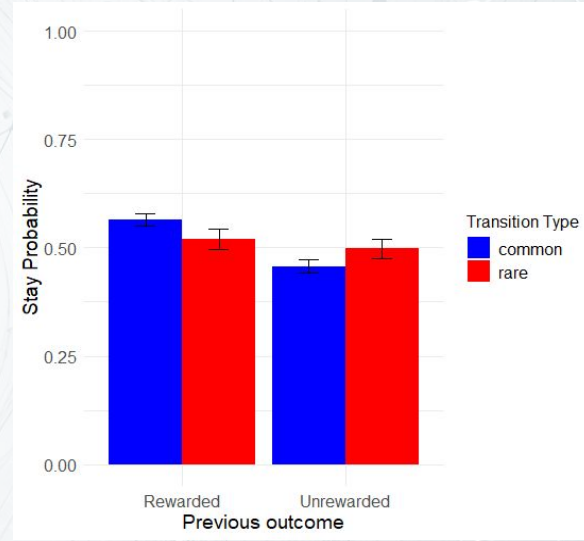


Easy



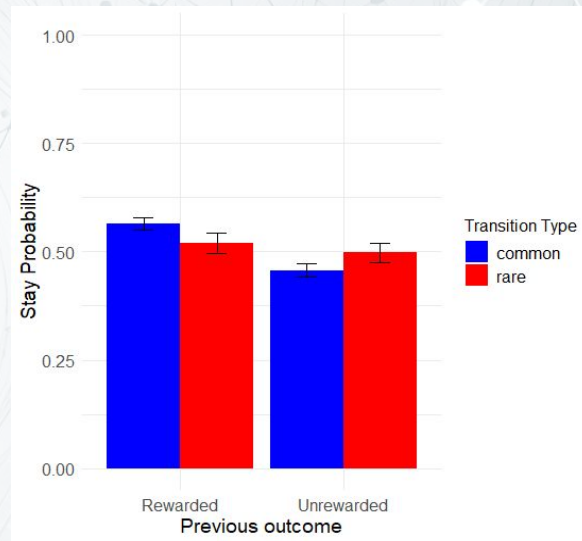

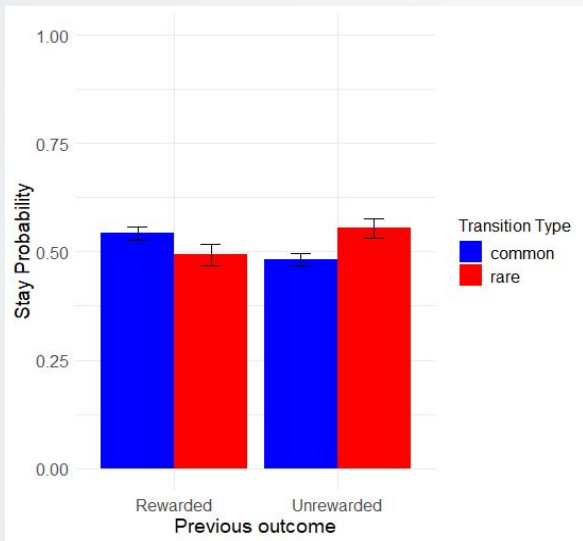





Where's $\omega=0$?




Where's $\omega=1$?




$\omega=0$




$\omega=0.5$



$\omega=1$



Grid search for omega (github)



Step-by-step estimation

[google sheet](#)



Thank you for listening :)

Extra reading:

1. See [Collins&Cockburn 2020](#) for a recent review on the topic.
2. An agent could also represent a model of the environment in terms of how outcomes are associated with one another (Moran et al., [2019](#), [2021a](#), [2021b](#)).
3. Recent results propose that a story-like framing of the instructions could lead participants to be much more model-based (Feher da Silva & Hare, [2020](#), [2023](#)).